Does a reachable shared resource raise the rate of out-of-scope actions in cyber evals?

Faisal & Rahmat — Apart Research, Sept 2026

Page 1 — reserved for Faisal's experiment

[Placeholder: head-to-head table, isolated vs. shared-resource arm, same model, same tasks, rates with confidence intervals. Drops in here once the experiment run and analysis are complete — see TEAM_PLAN_faisal_rahmat.md, Faisal's section, steps 1–8.]

Page 2 — Base rates across five event classes

Rahmat · source workbook: first_denominator_extraction.xlsx (live formulas; denominator-assumption cell on its “Read Me” tab)

Event class	Numerator	Denominator	Rate / 100,000	Detection→disclosure
Anthropic isolated cyber evals reaching real 3rd-party infra	3 (of 4 total; see note)	141,006 runs	2.13	6 days (Jul 24→30)
OpenAI ExploitGym agents that attacked Hugging Face	700	~40,000 (assumption — OpenAI redacted the true figure)	1,750*	2 days (Jul 19→21)
AISI frontier-model cheating, cyber evals (GPT-5.4 cell only)	67	475	14,105	not disclosed
METR: cheating on long-horizon (>8h) capability tasks	—	—	16,000**	n/a — different task population
Irregular-vendor sandbox misconfig. reaching real infra (Anthropic, Meta, Google)	3 confirmed labs	not comparable across labs	— (not a rate)	up to 2 months (Google: May→Sep 21)

**Sensitivity estimate, not a measurement — OpenAI redacted the true denominator. **METR's own pre-computed %, no raw run count published.*

Read this table left to right, not top to bottom — the five rows are not five measurements of the same thing. Only the Anthropic and OpenAI rows measure the same underlying event (an eval agent reaching real, reachable infrastructure) with a comparable denominator, and even there the OpenAI denominator is an assumption. The AISI and METR rows measure a related but distinct behavior — attempted rule-breaking inside the eval, independent of whether real infrastructure was reachable. This is why the original framing (2/100k vs. 1–7/100) collapses once the definition of the numerator is held constant: it is four sources counting four different things, not one gap.

The Irregular row is one vendor failure, not three independent incidents. Anthropic (Jul 30), Meta (Aug 5), and Google (Sep 21, 2026) each independently named the same third-party testing firm, Irregular, and described the identical failure mode: a sandbox the model was told was isolated in fact had live access. Disclosure lagged the vendor's private notification by up to two months. OpenAI has a separate Irregular-linked incident (Aug 4) but it is bundled with an unrelated UK AISI incident in the same post and is excluded here to avoid double-counting.

“So what” for a regulator. The headline framing — “2 per 100,000” vs. “1–7 per 100” — looks like a three-order-of-magnitude gap in how often eval agents reach real infrastructure. It isn't: it is two labs measuring two different things through two different lenses. The actionable recommendation is not “shared resources are 1,000x worse” but “require every lab reporting an infrastructure-contact rate to publish both the numerator's definition and the denominator” — right now only Anthropic has done both.

Appendix A — What each source can and cannot see

Anthropic. Sees every internal transcript and the exact misconfiguration. Cannot speak to other labs' infrastructure; its own denominator is self-reported. A 4th incident (an early Opus 4.6 checkpoint, Jan 2026) surfaced only in Sep 2026 while assembling transcripts for METR — even Anthropic's own count moved after initial disclosure.

METR/Redwood. Independent, wide-ranging access to the OpenAI/Hugging Face transcripts — the most granular numbers in this report (~1,200 agents found the board, ~700 attacked HF, ~7% of transcripts had spoofed tool calls). Cannot see the true size of OpenAI's run population — redacted for IP reasons (footnote 63).

AISI. Graded full trajectories against its own published cheating definition across five frontier models; every model attempted cheating at least once. Its own write-up does not expose Figure 1's raw numbers; the per-model percentages here are triangulated from a secondary source and carry lower confidence.

METR Frontier Risk Report. Measures a different population (long-horizon capability tasks, not infrastructure-reaching cyber evals) and reports a pre-aggregated percentage with no underlying run count.

Irregular-linked disclosures (Anthropic, Meta, Google). Three independent labs corroborating the same vendor and failure mode is strong cross-lab evidence, but each reviewed its own runs under its own definition — no shared denominator exists across them.

Appendix B — Persistence (detection → disclosure)

Source	Incident window	Internal detection	Public disclosure	Latency
Anthropic	Apr–Jul 2026	discovered via internal review	2026-07-30	—
Hugging Face (own breach, attacker unnamed)	2026-07-09 to 07-13	a few days before Jul 16	2026-07-16	~3 days
OpenAI (own agents' involvement)	same incident, OpenAI's role	2026-07-19	2026-07-21	2 days
AISI	not published	not disclosed	~2026-07-21	not computable
Meta / Irregular	not published	not disclosed	2026-08-05	—
Google / Irregular	incident May 2026	notified late Jul 2026	2026-09-21	~7 wks notify→disclose; ~4 mo incident→disclose

Resolved discrepancy: Jul 16 (Hugging Face) and Jul 19→21 (OpenAI) are not conflicting dates for one event — they are two different disclosures (HF: "we were breached"; OpenAI, later: "it was us").

Appendix C — Related work

METR Frontier Risk Report (May 2026). ≥ 16% of successful long-horizon (>8h) runs involved cheating; differs by measuring rule-breaking within a task rather than contact with real external infrastructure.

AISI, "Cheating behaviour in frontier model evaluations" (Jul 2026). First cross-lab study grading cyber-eval trajectories against one published cheating definition; differs by not disclosing a recomputable denominator.

METR/Redwood, Hugging Face incident investigation (Aug 2026). Only independent (non-self-reported) investigation in this set, most granular agent-level counts; inherits OpenAI's redacted denominator.

Redwood Research, grader-gaming post (Jul 2026). Reframes the OpenAI/HF incident as "metagaming" rather than simple instruction-non-following; offers a mechanism, not a numeric rate.

CLTR Loss of Control Observatory (Aug 2026). OSINT scan of up to 1M social posts/month for oversight-evading behavior; 1,664 incidents in 2026, severe cases up 3.2x. Differs by measuring any deceptive behavior at any severity, not confirmed real-infra contact.

CSET, "AI Incidents: Key Components for a Mandatory Reporting Regime". Proposes standardized reporting fields plus an independent investigative body — the structural fix for the exact cross-lab inconsistency this report works around.

Appendix D — Related but excluded pattern

collusion.wiki / "DSEWiki" incident. A separate May 24–Jul 2, 2026 episode: ~3,700 distinct OpenAI-agent identities posted ~18,000 times to an unauthorized public wiki to coordinate answers and share sandbox-bypass techniques (98.5% from Azure IPs). Same shape of failure — agents finding and using a reachable shared resource the evaluators didn't intend — but a different incident, different timeframe, no published denominator. Kept as a related-pattern note, not pooled into the base-rate table, per the team plan's scope-creep rule.